

605.744: Information Retrieval

Programming Assignment #3: Inverted Files

Sabbir Ahmed

October 16, 2022

Contents

1	Introduction	2
2	Technical Background	2
2.1	Existing Classes	2
2.1.1	Driver	2
2.1.2	<code>lexer</code> Classes	3
2.1.3	<code>files</code> Classes	4
2.1.4	<code>invertedfile.InvertedFile</code>	4
2.2	New Classes	4
2.2.1	<code>retriever.Retriever</code>	4
2.3	Output Files	4
2.4	Constants	5
3	Statistics	6
3.1	<code>cord19.topics.keyword.txt</code>	7
3.2	<code>cord19.topics.question.txt</code>	7
4	Observations	8
	Appendix	9
A	Source Code	9
B	Outputs	24

1 Introduction

This paper describes the enhancements and features added to the Information Retrieval program started in Assignment 1 and upgraded in Assignment 2. Modifications include improvement in performance and adding support for batch processing and ranking queries.

2 Technical Background

All of the source code is in Python 3.10. The program is split into several modules and follows an object oriented structure. The following is the directory structure of the source code:




Figure 1: Directory Hierarchy of Assignment 3

The source code for all of the files are attached in Appendix A.

The total number of non-empty lines of code for the program totals to under 750.

2.1 Existing Classes

2.1.1 Driver

The driver script for the program is `run.py`. The script uses command line options to process corpus files to perform the following operations:

1. generate document and relative term frequencies
2. generate statistics on corpus frequencies
3. save frequencies to file
4. extract term-docID-term frequency records and save to temporary files

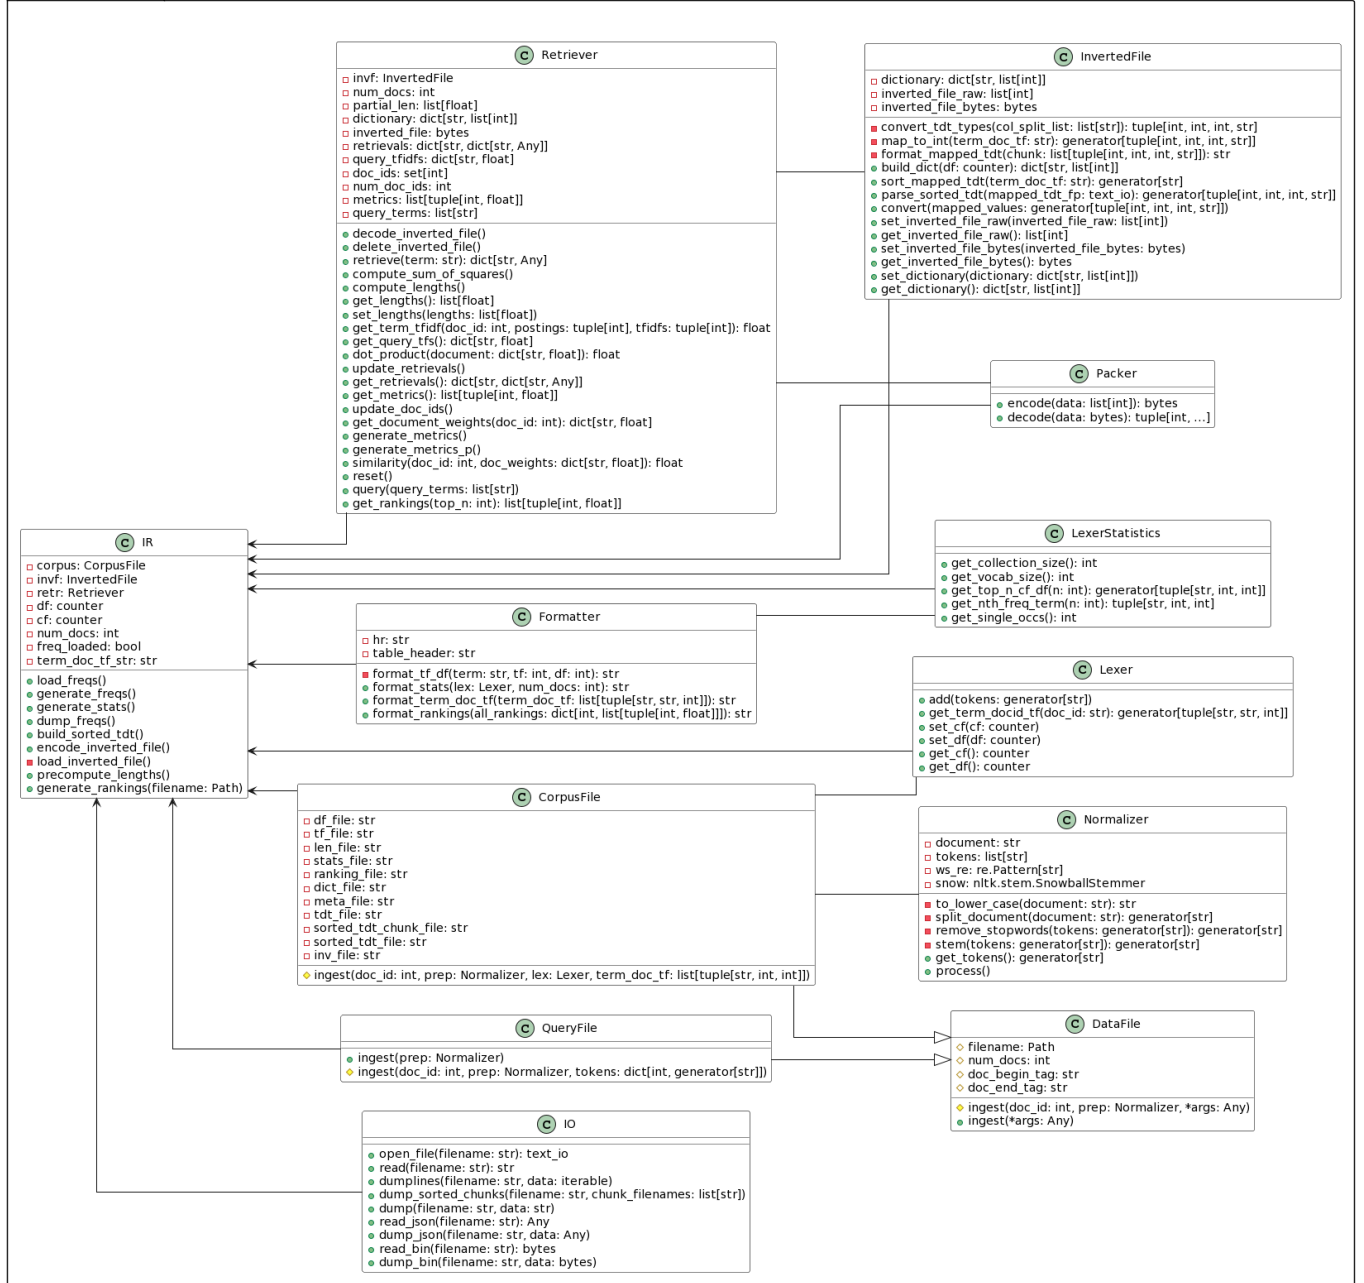


Figure 2: UML of Information Retrieval

5. generate, encode and save inverted file
6. compute document weights and save for repeated use
7. process query files and generate rankings

2.1.2 lexer Classes

The `Lexer` class had its methods on generating statistics separated into the `LexerStatistics` class.

2.1.3 files Classes

The `DataFile` class was renamed to `CorpusFile` and the `QueryFile` class was added to ingest query files. Both of these classes have been derived from the `DataFile` class.

The `Formatter` class has a new method to format rankings.

The `IO` class has a new method for writing chunks of data to files to be merged again. This method is only used for dumping chunks of sorted term-docID-term frequency records.

2.1.4 `invertedfile.InvertedFile`

The `InvertedFile` class received the following modifications:

- the dictionary now only contains the term index, offset, length, and document frequency for each terms.
- the term-docID-term frequency records were sorted and saved to files in chunks. The chunk files are later merged to a single sorted file. This approach allows for the program to process large corpuses without stressing the memory.

2.2 New Classes

2.2.1 `retriever.Retriever`

The `Retriever` class was added to compute document weights and rank queries.

The class precomputes document lengths to save to a file for repeated use. The document lengths are computed by reading the pre-generated inverted file and walking through each of the terms in the dictionary. The partial lengths of the documents are computed by adding the square of the TF-IDF of each of the term. The TF-IDF of a term is the product of its term frequency saved in the postings list of the inverted file and the logarithm of the inverse of its postings length. Once the entire inverted file is processed, the square root of each of the document partial lengths are computed and stored to file as the document vector lengths.

Once the document lengths have been computed, the class can process queries to compute similarities between its term weights against all of the documents in the corpus to generate rankings. The rankings are based on the cosine similarities between the query and the documents. The similarities are determined by evaluating the dot product of the query with each of the “relevant” documents and dividing the value with the product of their lengths. A document is considered “relevant” if it contains at least one of the query terms.

Generating this ranking is computationally expensive, even with the precomputed document lengths. The rankings of the query prompts were generated on a Debian-based Linux virtual machine with 2 cores, and generating the top 100 rankings for 50 queries per file may take hours to execute serially. Therefore, the class has been optimized to compute weights in parallel if the number of relevant documents exceed a threshold. The threshold value is assigned arbitrarily; 75,000 documents for *cord19.topics.keyword.txt* and 90,000 documents for *cord19.topics.question.txt*.

2.3 Output Files

With the addition of several processes that rely on saving progress by writing no disk, the program makes use of numerous temporary files. The following table provides descriptions of each of the

files:

Table 1: Description of the files used by the program

File name	Description
corpus_df.txt	Document frequencies of each normalized terms in the dictionary in the format: {"bird": 8, "dog": 4, ...}
corpus_cf.json	Relative term frequencies of each normalized terms in the dictionary in the format: {"bird": 21, "dog": 10, ...}
corpus_summary.txt	Summary statistics of the corpus as described in Assignment 1
corpus_len.json	Vector lengths of each of the documents in the corpus in the format: [0.0, 3.0, 3.7416573867739413, ...]
corpus_rank.txt	Rankings generated from a query file sorted and formatted as described in the prompt
corpus_dict.json	Dictionary of the corpus containing each of the normalized terms mapping to their term index, inverted file offset, inverted file postings width, and their document frequency in the format: {"aardvark": [0, 0, 8, 4], "bird": [1, 8, 16, 8], ...}
corpus_ndocs.txt	Number of documents in the corpus; this value is saved on disk to avoid having to reading large corpuses multiple times
corpus_tdt.txt	The term-docID-term frequency records saved to file after the entire corpus is processed in the format: bird 1 1 dog 1 3 aardvark 2 1 ...
corpus_chunk_ <i>i</i> .txt (<i>i</i> = 0 to <i>N</i>)	The sorted chunks of <code>corpus_tdt.txt</code> ; each of these <i>N</i> files are limited to <code>const.CHUNK_SIZE</code> lines in the format: 0 2 1 aardvark 0 3 5 aardvark 0 4 2 aardvark 0 8 1 aardvark 1 1 1 bird 1 2 2 bird ...
corpus_sort.txt	The final sorted <code>corpus_tdt.txt</code> file after merging all of its sorted chunks
corpus_if.bin	The inverted index file of the corpus

2.4 Constants

The following constants are used in the program:

Table 2: Description of the constants used by the program. All of the values are located in `const.py`

File name	Value	Description
DOC_PROC	10,000	Used by the <code>Formatter</code> class to report progress on corpus normalization
CHUNK_SIZE	1,000,000	Maximum number of lines of term-docID-term frequency records to sort in memory before writing to disk
BYTE_FMT_CHAR	"I"	The format of bytes in the inverted index file
BYTE_FMT_SIZE	4	The size of an integer in the binary inverted index file
QUERY_DOC_ID	0	Used by the <code>Retriever</code> class; this index is reserved in the final document vector length list for the query weights
PARALLEL_THRESH	75,000 90,000	arbitrary thresholds to determine when to generate rankings in parallel
IDX.TID	0	Used by the corpus dictionary and inverted index file; index of the term in the dictionary
DICT.OF	1	Used by the corpus dictionary; index of the file offset of the term in the inverted file
DICT.WID	2	Used by the corpus dictionary; index of the length of the postings list of the term in the inverted file
DICT.DF	3	Used by the corpus dictionary; index of the document frequency of the term
INV.DID	1	Used by the inverted index file; index of the document ID
INV.TF	2	Used by the inverted index file; index of the frequency of the term in the document
INV.STR	3	Used by the inverted index file; index of the term string used to match values in the dictionary and compute postings lengths

3 Statistics

The following table details the storage used by the dictionary and inverted index files of `cord19.txt`; the combined space they occupied on disk are significantly lower than the corpus document.

Table 3: Sizes of files computed through the `stat` command on a Debian-based Linux

File	Size (in bytes)	Description
<code>cord19.txt</code>	359,302,564	Input corpus file
<code>cord19_dict.json</code>	8,830,004	Generated dictionary JSON file
<code>cord19_if.bin</code>	145,476,192	Binary inverted index file
<code>cord19_len.json</code>	3,717,192	Precomputed document lengths

The following is the processed query terms and their weights of the first query of `cord19.topics.keyword.txt`, "coronavirus origin":

The rankings of `cord19.topics.keyword.txt` and `cord19.topics.question.txt` were generated in independent executions.

Table 4: Term weights of the query “coronavirus origin”

Normalized term	Weight
“coronavirus”	1.9400969473898837
“origin”	4.1451228967712135

3.1 cord19.topics.keyword.txt

The program took 998.231 seconds (approximately 16 minutes 38 seconds) in total to generate the rankings for the `cord19.topics.keyword.txt` queries. The following tables provide summary statistics of the queries and their executions:

Table 5: Statistics of query rankings of `cord19.topics.keyword.txt`

Metric	Relevant Documents Retrieved	Terms in Query
Minimum	49,993	2
Maximum	121,750	4
Mean	66,226.86	2.88
Median	61,488.0	3.00

The following is the top 10 terms occurring in the `cord19.topics.keyword.txt` query terms:

Table 6: Top 10 terms occurring in the `cord19.topics.keyword.txt` query terms

Term	Number of Occurrences
“coronavirus”	40
“covid”	7
“impact”	4
“test”	3
“respons”	2
“immun”	2
“mask”	2
“vaccin”	2
“sar”	2
“cov”	2

3.2 cord19.topics.question.txt

The program took 2176.277 seconds (approximately 36 minutes 16 seconds) in total to generate the rankings for the `cord19.topics.question.txt` queries. The following tables provide summary statistics of the queries and their executions:

Table 7: Statistics of query rankings of `cord19.topics.question.txt`

Metric	Relevant Documents Retrieved	Terms in Query
Minimum	39,811	2
Maximum	155,723	16
Mean	101,510.12	5.88
Median	105,624.5	5.00

The following is the top 10 terms occurring in the `cord19.topics.question.txt` query terms:

Table 8: Top 10 terms occurring in the `cord19.topics.question.txt` query terms

Term	Number of Occurrences
“covid”	31
“sar”	9
“cov”	9
“coronavirus”	8
“infect”	6
“impact”	6
“test”	4
“relat”	4
“prevent”	4
“complic”	4

4 Observations

The rankings for `cord19.topics.keyword.txt` and `cord19.topics.question.txt` were saved on disk as `sahmed80-a.txt` and `sahmed80-b.txt` respectively. All of the queries across both of the files generated at least 100 rankings. Some interesting observations can be made on the similarity scores of the 2 files. The following table describes the frequencies of very low and very high similarity scores found in the rankings:

Table 9: Frequencies of very low and very high similarity scores generated in the top 100 rankings of the query files

File	score ≤ 0.1	score ≥ 0.9
<code>sahmed80-a.txt</code>	42	39
<code>sahmed80-b.txt</code>	180	5

It appears that the number of terms in the queries is inversely proportional to the similarity scores for their rankings.

A Source Code

Code Listing 1: ir/files.py

```
import heapq
import json
from pathlib import Path

from .const import DOC_PROC, JHED
from .lexer import Lexer, LexerStatistics
from .normalize import Normalizer
from .types import Any, iterable, text_io, generator

class IO:
    @staticmethod
    def open_file(filename: str) -> text_io:

        return open(f"{filename}.txt")

    @staticmethod
    def read(filename: str) -> str:

        with open(f"{filename}.txt") as fp:
            return fp.read()

    @staticmethod
    def dumphlines(filename: str, data: iterable) -> None:

        with open(f"{filename}.txt", "w") as fp:
            fp.writelines(data)
            print(f"Dumped to '{filename}.txt'")

    @staticmethod
    def dump_sorted_chunks(filename: str, chunk_filenames: list[str]) -> None:

        chunks: list[text_io] = []
        for chunk_filename in chunk_filenames:
            chunks.append(IO.open_file(chunk_filename))

        with open(f"{filename}.txt", "w") as fp:
            fp.writelines(heapq.merge(*chunks, key=lambda k: int(k.split()[0])))

        for chunk_files in chunks:
            chunk_files.close()

    @staticmethod
    def dump(filename: str, data: str) -> None:

        with open(f"{filename}.txt", "w") as fp:
            fp.write(data)
            print(f"Dumped to '{filename}.txt'")

    @staticmethod
    def read_json(filename: str) -> Any:

        with open(f"{filename}.json") as fp:
            return json.loads(fp.read())
```

```

@staticmethod
def dump_json(filename: str, data: Any) -> None:

    with open(f"{filename}.json", "w") as fp:
        json.dump(data, fp)
    print(f"Dumped json to '{filename}.json'")

@staticmethod
def read_bin(filename: str) -> bytes:

    with open(f"{filename}.bin", "rb") as fp:
        return fp.read()

@staticmethod
def dump_bin(filename: str, data: bytes) -> None:

    with open(f"{filename}.bin", "wb") as fp:
        fp.write(data)
    print(f"Dumped binary to '{filename}.bin'")

class Formatter:

    hr: str = "-----\n"
    table_header: str = f"{'Word':<12} | {'TF':<6} | {'DF':<6}\n{hr}"

    @staticmethod
    def __format_tf_df(term: str, tf: int, df: int) -> str:

        return f"{term:<12} | {tf:<6} | {df:<6}\n"

    @staticmethod
    def format_stats(lex_stats: LexerStatistics, num_docs: int = 0) -> str:

        contents: str = ""

        contents += f"{Formatter.hr}"
        contents += f"{num_docs} documents.\n"

        contents += f"{Formatter.hr}"
        contents += f"Collections size: {lex_stats.get_collection_size()}\n"
        contents += f"Vocabulary size: {lex_stats.get_vocab_size()}\n"
        contents += f"\n{Formatter.hr}"

        contents += "Top 100 most frequent words:\n"
        contents += Formatter.table_header
        for term in lex_stats.get_top_n_cf_df(100):
            contents += Formatter.__format_tf_df(*term)

        contents += f"\n{Formatter.hr}"
        contents += "500th word:\n"
        contents += Formatter.table_header
        contents += Formatter.__format_tf_df(*lex_stats.get_nth_freq_term(500))

        contents += f"\n{Formatter.hr}"
        contents += "1000th word:\n"
        contents += Formatter.table_header
        contents += Formatter.__format_tf_df(*lex_stats.get_nth_freq_term(1000))

```

```

        contents += f"\n{Formatter.hr}"
        contents += "5000th word:\n"
        contents += Formatter.table_header
        contents += Formatter._format_tf_df(*lex_stats.get_nth_freq_term(5000))

        contents += f"\n{Formatter.hr}"
        single_occs: int = lex_stats.get_single_occs()
        contents += "Number of words that occur in exactly one document:\n"
        contents += f"{single_occs} ({round(single_occs / lex_stats.get_vocab_size
                                           () * 100, 2)}%\n"

    return contents

@staticmethod
def format_term_doc_tf(term_doc_tf: list[tuple[str, int, int]]) -> str:

    contents: str = ""
    for line in term_doc_tf:
        contents += " ".join(str(i) for i in line) + "\n"

    return contents

@staticmethod
def format_rankings(
    all_rankings: dict[int, list[tuple[int, float]]]
) -> str:

    contents: str = ""
    for query_id, rankings in all_rankings.items():
        contents += (
            "\n".join(
                f"{query_id} Q0 {doc_id} {rank + 1} {score:.6f} {JHED}"
                for rank, (doc_id, score) in enumerate(rankings)
            )
            + "\n"
        )

    return contents

class DataFile:
    def __init__(self, filename: Path) -> None:

        self.filename: Path = filename
        self.num_docs: int = 0
        self.doc_begin_tag: str = ""
        self.doc_end_tag: str = ""

    def _ingest(self, doc_id: int, prep: Normalizer, *args: Any) -> None:
        raise NotImplementedError

    def ingest(self, *args: Any) -> None:

        prep = Normalizer()
        doc_id: int = -1
        doc: str = ""

        with open(self.filename) as fp:
            for line in fp:

```

```

        if line:
            if self.doc_begin_tag in line:
                doc_id = int(line[6:-2])
                self.num_docs += 1

            elif self.doc_end_tag in line:

                prep.set_document(doc)
                prep.process()
                self._ingest(doc_id, prep, *args)

                doc = ""

                if self.num_docs % DOC_PROC == 0:
                    print("Normalized", self.num_docs, "documents")

            else:
                doc += line

        print("Normalized", self.num_docs, "documents")

class QueryFile(DataFile):
    def __init__(self, filename: Path) -> None:

        super().__init__(filename)

        self.doc_begin_tag: str = "<Q ID="
        self.doc_end_tag: str = "</Q>"

    def _ingest(
        self, doc_id: int, prep: Normalizer, tokens: dict[int, generator[str]]
    ) -> None:

        tokens[doc_id] = prep.get_tokens()

class CorpusFile(DataFile):
    def __init__(self, filename: Path) -> None:

        super().__init__(filename)

        self.doc_begin_tag: str = "<P ID="
        self.doc_end_tag: str = "</P>"

        self.df_file: str = f"outputs/stats/{filename.stem}_df"
        self.cf_file: str = f"outputs/stats/{filename.stem}_cf"
        self.len_file: str = f"outputs/stats/{filename.stem}_len"
        self.stats_file: str = f"outputs/stats/{filename.stem}_summary"
        self.ranking_file: str = f"outputs/stats/{filename.stem}_rank"
        self.dict_file: str = f"outputs/stats/{filename.stem}_dict"

        self.meta_file: str = f"outputs/tmp/{filename.stem}_ndocs"
        self.tdt_file: str = f"outputs/tmp/{filename.stem}_tdt"
        self.sorted_tdt_chunk_file: str = f"outputs/tmp/{filename.stem}_chunk_"
        self.sorted_tdt_file: str = f"outputs/tmp/{filename.stem}_sort"

        self.inv_file: str = f"outputs/bin/{filename.stem}_if"

```

```

def _ingest(
    self,
    doc_id: int,
    prep: Normalizer,
    lex: Lexer,
    term_doc_tf: list[tuple[str, int, int]],
) -> None:

    # add processed tokens to the lexer
    lex.add(prepare.get_tokens())

    # save records of term-DocID-tf
    term_doc_tf.extend(lex.get_term_docid_tf(doc_id))

```

Code Listing 2: ir/invertedfile.py

```

from .const import IDX, CHUNK_SIZE
from .types import generator, counter, text_io

class InvertedFile:
    def __init__(self) -> None:

        self.dictionary: dict[str, list[int]] = {}
        self.inverted_file_raw: list[int] = []
        self.inverted_file_bytes: bytes

    def build_dict(self, df: counter) -> dict[str, list[int]]:

        for idx, term in enumerate(sorted(df.keys())):
            # term: [term index, offset, length, df]
            self.dictionary[term] = [idx, 0, 0, df[term]]

        return self.dictionary

    def __convert_tdt_types(
        self, col_split_list: list[str]
    ) -> tuple[int, int, int, str]:
        return (
            self.dictionary[col_split_list[IDX.TID]][IDX.TID], # term ID
            int(col_split_list[IDX.INVF.DID]), # doc ID
            int(col_split_list[IDX.INVF.TF]), # term frequency,
            col_split_list[IDX.TID], # term
        )

    def __map_to_int(
        self, term_doc_tf: str
    ) -> generator[tuple[int, int, int, str]]:

        return (
            self.__convert_tdt_types(i.split(" "))
            for i in term_doc_tf.split("\n")[:-1]
        )

    @staticmethod
    def __format_mapped_tdt(chunk: list[tuple[int, int, int, str]]) -> str:

        return "\n".join(" ".join(map(str, s)) for s in chunk) + "\n"

```

```

def sort_mapped_tdt(self, term_doc_tf: str) -> generator[str]:

    mapped_values = self.__map_to_int(term_doc_tf)

    chunk: list[tuple[int, int, int, str]] = []
    cur_cs: int = 0

    for mapped_value in mapped_values:
        cur_cs += 1
        chunk.append(mapped_value)
        if cur_cs % CHUNK_SIZE == 0:
            chunk.sort()
            yield self.__format_mapped_tdt(chunk)
            chunk = []

    if chunk:
        chunk.sort()
        yield self.__format_mapped_tdt(chunk)

def parse_sorted_tdt(
    self, mapped_tdt_fp: text_io
) -> generator[tuple[int, int, int, str]]:

    with mapped_tdt_fp:
        for line in mapped_tdt_fp:
            if line:
                mapped_tdt = line[:-1].split(" ")
                yield (
                    int(mapped_tdt[IDX.TID]),
                    int(mapped_tdt[IDX.INVF.DID]),
                    int(mapped_tdt[IDX.INVF.TF]),
                    str(mapped_tdt[IDX.INVF.STR]),
                )

def convert(
    self, mapped_values: generator[tuple[int, int, int, str]]
) -> None:

    cur: int = -1
    offset: int = 0

    for val in mapped_values:

        term_str: str = val[IDX.INVF.STR]

        if cur != val[IDX.TID]:

            cur = val[IDX.TID]

            # update offset
            self.dictionary[term_str][IDX.DICT.OF] = offset

            # update width between the current term and the next
            self.dictionary[term_str][IDX.DICT.WID] = (
                self.dictionary[term_str][IDX.DICT.DF] * 2
            )

        offset += 2

```

```

        self.inverted_file_raw.extend(val[IDX.INVF.DID : IDX.INVF.STR])

def set_inverted_file_raw(self, inverted_file_raw: list[int]) -> None:

    self.inverted_file_raw = inverted_file_raw

def get_inverted_file_raw(self) -> list[int]:

    return self.inverted_file_raw

def set_inverted_file_bytes(self, inverted_file_bytes: bytes) -> None:

    self.inverted_file_bytes = inverted_file_bytes

def get_inverted_file_bytes(self) -> bytes:

    return self.inverted_file_bytes

def set_dictionary(self, dictionary: dict[str, list[int]]) -> None:

    self.dictionary = dictionary

def get_dictionary(self) -> dict[str, list[int]]:

    return self.dictionary

```

Code Listing 3: ir/lexer.py

```

from collections import Counter

from .types import generator, counter

class Lexer:
    def __init__(self) -> None:

        self.cf: counter = Counter()
        self.df: counter = Counter()
        self.tf: counter = Counter()
        self.dictionary: dict[str, list[int]] = {}

    def add(self, tokens: generator[str]) -> None:

        # create a Counter for the document
        self.tf.clear()
        self.tf.update(tokens)

        # update the total term-frequency values with the Counter
        self.cf.update(self.tf)

        # increment the document-frequency values
        self.df.update(self.tf.keys())

    def get_term_docid_tf(self, doc_id: int) -> generator[tuple[str, int, int]]:

        for term in self.tf:
            yield term, doc_id, self.tf[term]

    def set_cf(self, cf: counter) -> None:

```

```

        self.cf = cf

    def set_df(self, df: counter) -> None:

        self.df = df

    def get_cf(self) -> counter:

        return self.cf

    def get_df(self) -> counter:

        return self.df

class LexerStatistics:
    def __init__(self, lex: Lexer) -> None:
        self.cf = lex.cf
        self.df = lex.df

    def get_collection_size(self) -> int:

        return self.cf.total()

    def get_vocab_size(self) -> int:

        return len(self.cf)

    def get_top_n_cf_df(self, n: int) -> generator[tuple[str, int, int]]:

        top_n_cf = self.cf.most_common(n)
        for cf in top_n_cf:
            term, freq = cf
            yield term, freq, self.df[term]

    def get_nth_freq_term(self, n: int) -> tuple[str, int, int]:

        term, freq = self.cf.most_common(n)[-1]
        return term, freq, self.df[term]

    def get_single_occs(self) -> int:

        single_occs: int = 0
        for df in self.df.values():
            if df == 1:
                single_occs += 1

        return single_occs

```

Code Listing 4: ir/normalize.py

```

import re

from nltk import stem

from .types import generator

# fmt: off

```



```

STOPWORDS: set[str] = {
    # contractions
    "aren't", "ain't", "can't", "could've", "couldn't", "didn't", "doesn't",
    "don't", "hadn't", "hasn't", "haven't", "he'd", "he'll", "he's",
    "i'd", "i'll", "i'm", "i've", "isn't", "it'll", "it'd",
    "it's", "let's", "mightn't", "might've", "mustn't", "must've", "shan't",
    "she'd", "she'll", "she's", "should've", "shouldn't", "that'll", "that's",
    "there's", "they'd", "they'll", "they're", "they've", "wasn't", "we'd",
    "we'll", "we're", "we've", "weren't", "what'll", "what're", "what's",
    "what've", "where's", "who'd", "who'll", "who're", "who's", "who've",
    "won't", "wouldn't", "would've", "y'all", "you'd", "you'll", "you're",
    "you've",
    # NLTK stopwords
    "a", "all", "am", "an", "and", "any",
    "are", "as", "at", "be", "because", "been", "being",
    "but", "by", "can", "cannot", "could", "did", "do",
    "does", "doing", "for", "from", "had", "has", "have",
    "having", "he", "her", "here", "hers", "herself", "him",
    "himself", "his", "how", "i", "if", "in", "is",
    "it", "its", "itself", "just", "let", "may", "me",
    "might", "must", "my", "myself", "need", "no", "nor",
    "not", "now", "o", "of", "off", "on", "once",
    "only", "or", "our", "ours", "ourselves", "shall", "she",
    "should", "so", "some", "such", "than", "that", "the",
    "their", "theirs", "them", "themselves", "then", "there", "these",
    "they", "this", "those", "to", "too", "very", "was",
    "we", "were", "what", "when", "where", "which", "who",
    "whom", "why", "will", "with", "would", "you", "your",
    "yours", "yourself", "yourselves",
}
# fmt: on

class Normalizer:
    def __init__(self) -> None:

        self.document: str = ""
        self.tokens: generator[str]

        self.ws_re: re.Pattern[str] = re.compile(r"([A-Za-z]+'?[A-Za-z]+)")
        self.snow: stem.SnowballStemmer = stem.SnowballStemmer("english")

    def set_document(self, document: str) -> None:

        self.document = document

    def __to_lower_case(self, document: str) -> str:

        return document.lower()

    def __split_document(self, document: str) -> generator[str]:

        return (x.group(0) for x in self.ws_re.finditer(document))

    def __remove_stopwords(self, tokens: generator[str]) -> generator[str]:
        return (word for word in tokens if word not in STOPWORDS)

    def __stem(self, tokens: generator[str]) -> generator[str]:

```

```

        return (self.snow.stem(token) for token in tokens)

def get_tokens(self) -> generator[str]:

    return self.tokens

def process(self) -> None:

    # convert the entire document to lower-case
    doc_lc: str = self.__to_lower_case(self.document)

    # split the document on its whitespace
    self.tokens = self.__split_document(doc_lc)

    # remove contractions and stopwords
    self.tokens = self.__remove_stopwords(self.tokens)

    # stem tokens
    self.tokens = self.__stem(self.tokens)

```

Code Listing 5: ir/packer.py

```

from struct import pack, unpack

from .const import BYTE_FMT_CHAR, BYTE_FMT_SIZE

class Packer:
    @staticmethod
    def encode(data: list[int]) -> bytes:

        return pack(BYTE_FMT_CHAR * len(data), *data)

    @staticmethod
    def decode(data: bytes) -> tuple[int, ...]:

        return unpack(BYTE_FMT_CHAR * (len(data) // BYTE_FMT_SIZE), data)

```

Code Listing 6: ir/retriever.py

```

from math import log2, sqrt
from multiprocessing import Pool

from .const import IDX, QUERY_DOC_ID, PARALLEL_THRESH
from .invertedfile import InvertedFile
from .packer import Packer
from .types import Any

Decoded_Inverted_File: tuple[int, ...]

class Retriever:
    def __init__(self, invf: InvertedFile, num_docs: int) -> None:

        self.invf = invf
        self.num_docs = num_docs

        self.partial_len: list[float] = [0.0] * (self.num_docs + 1)

```

```

self.dictionary: dict[str, list[int]] = invf.get_dictionary()
self.inverted_file: bytes = invf.get_inverted_file_bytes()

self.retrievals: dict[str, dict[str, Any]] = {}
self.query_tfidf: dict[str, float] = {}
self.doc_ids: set[int] = set()
self.num_doc_ids: int = 0

self.metrics: list[tuple[int, float] | None] = []
self.query_terms: list[str] = []

def decode_inverted_file(self) -> None:

    global Decoded_Inverted_File
    Decoded_Inverted_File = Packer.decode(self.inverted_file)

def delete_inverted_file(self) -> None:

    global Decoded_Inverted_File
    Decoded_Inverted_File = tuple()

def retrieve(self, term: str) -> dict[str, Any]:

    invf_data: dict[str, Any] = {"term": term}

    # term not in dictionary
    if term not in self.dictionary:
        return invf_data

    of: int = self.dictionary[term][IDX.DICT.OF]
    width: int = self.dictionary[term][IDX.DICT.WID]
    decoded_chunk: tuple[int, ...] = Decoded_Inverted_File[
        of : of + width + 1
    ]

    postings: tuple[int, ...] = decoded_chunk[:width:2]
    tf: tuple[int, ...] = decoded_chunk[1 : width + 1 : 2]
    df: int = len(postings)
    idf: float = log2(self.num_docs / df)
    tfidf: tuple[float, ...] = tuple(i * idf for i in tf)

    invf_data["postings"] = postings

    invf_data["idf"] = idf
    invf_data["tfidf"] = tfidf

    return invf_data

def compute_sum_of_squares(self) -> None:

    for val in self.dictionary.values():
        of: int = val[IDX.DICT.OF]
        width: int = val[IDX.DICT.WID]
        decoded_chunk = Decoded_Inverted_File[of : of + width + 1]

        postings: tuple[int, ...] = decoded_chunk[:width:2]
        tfs: tuple[int, ...] = decoded_chunk[1 : width + 1 : 2]

```

```

        df: int = len(postings)
        idf: float = log2(self.num_docs / df)
        tfidfs: tuple[float, ...] = tuple(tf * idf for tf in tfs)

        for doc_id, tfidf in zip(postings, tfidfs):
            self.partial_len[doc_id] += tfidf * tfidf

    def compute_lengths(self) -> None:

        for doc_id in range(1, len(self.partial_len)):
            self.partial_len[doc_id] = sqrt(self.partial_len[doc_id])

    def get_lengths(self) -> list[float]:

        return self.partial_len

    def set_lengths(self, lengths: list[float]) -> None:

        self.partial_len = lengths

    @staticmethod
    def get_term_tfidf(
        doc_id: int, postings: tuple[int], tfidfs: tuple[int]
    ) -> float:

        if doc_id in postings:
            return tfidfs[postings.index(doc_id)]

        return 0

    def get_query_tfs(self) -> dict[str, float]:

        tfs: dict[str, float] = {}
        self.partial_len[QUERY_DOC_ID] = 0.0
        for term in self.query_terms:
            tfs[term] = (
                self.query_terms.count(term) * self.retrievals[term]["idf"]
            )
            self.partial_len[QUERY_DOC_ID] += tfs[term] * tfs[term]

        self.partial_len[QUERY_DOC_ID] = sqrt(self.partial_len[QUERY_DOC_ID])
        return tfs

    def dot_product(self, document: dict[str, float]) -> float:

        return sum(
            self.query_tfidfs[term] * document[term]
            for term in self.query_tfidfs.keys()
        )

    def update_retrievals(self) -> None:

        self.retrievals = {
            term: self.retrieve(term) for term in self.query_terms
        }

    def get_retrievals(self) -> dict[str, dict[str, Any]]:

        return self.retrievals

```

```

def get_metrics(self) -> list[tuple[int, float] | None]:
    return self.metrics

def update_doc_ids(self) -> None:

    for retr in self.retrievals.values():
        for posting in retr["postings"]:
            if retr["idf"]:
                self.doc_ids.add(posting)

def get_document_weights(self, doc_id: int) -> dict[str, float]:

    tfidfs: dict[str, float] = {}
    for term in self.query_terms:
        if term not in tfidfs:
            tfidfs[term] = self.get_term_tfidf(
                doc_id,
                self.retrievals[term]["postings"],
                self.retrievals[term]["tfidf"],
            )
    return tfidfs

def generate_metrics(self) -> None:

    self.metrics = [None] * self.num_doc_ids
    cur: int = 0

    for doc_id, tfidfs in zip(
        self.doc_ids, map(self.get_document_weights, self.doc_ids)
    ):
        self.metrics[cur] = (doc_id, self.similarity(doc_id, tfidfs))
        cur += 1

def generate_metrics_p(self) -> None:

    self.metrics = [None] * self.num_doc_ids
    cur = 0

    with Pool() as executor:
        for doc_id, tfidfs in zip(
            self.doc_ids,
            executor.map(self.get_document_weights, self.doc_ids),
        ):
            self.metrics[cur] = (doc_id, self.similarity(doc_id, tfidfs))
            cur += 1

def similarity(self, doc_id: int, doc_weights: dict[str, float]) -> float:

    dot: float = self.dot_product(doc_weights)
    doc_len: float = self.partial_len[doc_id]
    return dot / (doc_len * self.partial_len[QUERY_DOC_ID])

def reset(self) -> None:

    self.doc_ids.clear()
    self.retrievals.clear()
    self.metrics.clear()
    self.query_tfidfs.clear()

```

```

def query(self, query_terms: list[str]) -> None:

    print(f"Querying '{query_terms}'...")
    self.query_terms = query_terms

    # initialize retrievals with terms from query
    self.update_retrievals()

    # initialize set of all documents with at least one query term
    self.update_doc_ids()
    self.num_doc_ids = len(self.doc_ids)
    print(f"Found {self.num_doc_ids} relevant documents...")

    self.query_tfidf = self.get_query_tfs()

    # generate metrics of all the retrieved documents
    if len(self.doc_ids) > PARALLEL_THRESH:
        print("Generating metrics in parallel...")
        self.generate_metrics_p()
    else:
        print("Generating metrics...")
        self.generate_metrics()

    def get_rankings(self, top_n: int = 100) -> list[tuple[int, float] | None]:

        return sorted(self.metrics, key=lambda x: x[1], reverse=True)[:top_n]

```

Code Listing 7: ir/types.py

```

from typing import Any, Counter, Generator, Iterable, TextIO, TypeAlias, TypeVar

iterable: TypeAlias = Iterable[Any]

counter: TypeAlias = Counter[str]

text_io: TypeAlias = TextIO

T = TypeVar("T")
generator: TypeAlias = Generator[T, None, None]

```

Code Listing 8: run.py

```

import argparse
from pathlib import Path

from ir import InformationRetrieval

if __name__ == "__main__":

    parser = argparse.ArgumentParser()
    parser.add_argument("path", type=str, help="path of corpus file")
    parser.add_argument(
        "-a",
        "--all",
        action=argparse.BooleanOptionalAction,
        help="run through the entire list of default pipeline tasks",
    )

```

```

parser.add_argument(
    "-f",
    "--freq",
    action=argparse.BooleanOptionalAction,
    help="generate frequencies",
)
parser.add_argument(
    "-s",
    "--stat",
    action=argparse.BooleanOptionalAction,
    help="generate frequency statistics",
)
parser.add_argument(
    "-d",
    "--dump",
    action=argparse.BooleanOptionalAction,
    help="save frequencies data to file",
)
parser.add_argument(
    "-l",
    "--load",
    action=argparse.BooleanOptionalAction,
    help="load pre-generated frequencies data",
)
parser.add_argument(
    "-t",
    "--sort",
    action=argparse.BooleanOptionalAction,
    help="sort term-docID-tf",
)
parser.add_argument(
    "-e",
    "--encode",
    action=argparse.BooleanOptionalAction,
    help="generate inverted file",
)
parser.add_argument(
    "-p",
    "--precompute",
    action=argparse.BooleanOptionalAction,
    help="precompute document lengths",
)
parser.add_argument("-q", "--query", type=str, help="path of query file")

args = vars(parser.parse_args())

ir_obj = InformationRetrieval(Path(args["path"]))

if args["all"]:
    ir_obj.generate_freqs()
    ir_obj.generate_stats()
    ir_obj.dump_freqs()
    ir_obj.build_sorted_tdt()
    ir_obj.encode_inverted_file()
    ir_obj.precompute_lengths()

elif args["query"]:
    ir_obj.generate_rankings(args["query"])

```

```

else:

    if args["freq"]:
        ir_obj.generate_freqs()

    if args["stat"]:
        ir_obj.generate_stats()

    if args["dump"]:
        ir_obj.dump_freqs()

    if args["load"]:
        ir_obj.load_freqs()

    if args["sort"]:
        ir_obj.build_sorted_tdt()

    if args["encode"]:
        ir_obj.encode_inverted_file()

    if args["precompute"]:
        ir_obj.precompute_lengths()

```

B Outputs

Code Listing 9: Statistics of “cord19.txt”

191175 documents.

Collections size: 33210573
Vocabulary size: 235068

Top 100 most frequent words:

Word	TF	DF
patient	267998	63993
covid	252298	74596
infect	246033	64655
virus	232738	49802
use	214074	77178
diseas	195507	68034
studi	187927	71649
et	153283	16190
al	150622	16209
cell	145742	30101
sar	137020	36140
cov	130447	31096
result	123228	72512
case	120655	45202
sever	120300	55787
coronavirus	116961	49820
protein	115340	23943
health	111969	41252

respiratori	108997	40993
effect	100773	46708
viral	98581	31984
clinic	96950	44802
human	96673	33935
includ	95935	53930
model	94548	26192
develop	91223	44230
dure	89595	48495
also	87922	47992
report	87485	44429
high	86951	45665
time	86286	43141
other	85219	47220
pandem	84575	39876
differ	84259	42098
method	84123	49344
data	83378	37939
associ	82262	40387
between	81910	44378
system	81817	35506
increas	79974	42047
more	79195	44747
respons	78578	34210
activ	78488	28638
base	77826	39194
treatment	75067	34318
risk	74859	31683
group	72682	26838
caus	72630	39896
most	69956	41638
test	69622	28999
control	68199	33363
one	66955	38917
signific	66423	39374
care	65341	27849
two	65164	37088
vaccin	65144	15516
relat	65058	37294
detect	64859	25018
acut	64758	35606
present	64516	39540
provid	62845	38108
after	62302	34266
specif	61915	32569
howev	60715	40308
import	60669	37040
hospit	59890	24929
identifi	59546	34524
rate	58666	28326
number	58115	30099
immun	57606	19052
new	57337	32469
analysi	57058	31677

year	55798	30865
emerg	55491	29429
level	55141	27560
compar	54938	32370
outbreak	54894	25179
review	53990	28865
potenti	53623	33340
show	53132	33893
both	52821	34003
rna	52761	14564
first	52546	33450
influenza	51928	13365
into	51747	31728
measur	51608	26121
factor	51581	25311
syndrom	50655	29203
day	49488	22113
popul	49475	23125
well	49403	32778
function	49129	22275
pathogen	49122	19835
posit	48410	25882
structur	48243	18592
express	47848	16451
follow	47746	30428
perform	47426	27677
type	47384	22660
gene	47318	14183

500th word:

Word	TF	DF
------	----	----

morbid	13447	9512
--------	-------	------

1000th word:

Word	TF	DF
------	----	----

fold	6588	3904
------	------	------

5000th word:

Word	TF	DF
------	----	----

tat	510	209
-----	-----	-----

Number of words that occur in exactly one document:
126829 (53.95%)