

605.744: Information Retrieval

Final Exam

Sabbir Ahmed

December 4, 2022

1. (a) Suppose the document collection is very large, and that the index will not fit in the available RAM. Describe an indexing algorithm that works when memory is small compared to the size of the index.
(b) Once an inverted file has been created, it is possible to calculate document vector lengths for a TF/IDF cosine model. This pre-calculation makes query-time performance much more efficient. Explain how right after creating the inverted file document vector lengths can be efficiently computed for all documents in parallel using one traversal (i.e., one single pass) over the inverted file.

Table 1: Cardinalities of set-difference sets with various n-grams and normalization parameters

N-Gram	Normalized	$ G - O $	$ O - G $
6	True	10	12
6	False	10	12
5	True	3	6
5	False	4	9
4	True	3	6
4	False	3	6
3	True	2	4
3	False	2	4
2	True	1	2
2	False	3	2
1	True	2	1
1	False	2	2

2. The three major problems in text retrieval are: (a) polysemy; (b) synonymy; and, (c) morphology. Briefly explain each issue and how it can lower performance. Give an example of each phenomena.

Answer:

- (a) Polysemy refers to words that can have multiple meanings depending on the context. For example, *space* can refer to its noun version of unoccupied area. The unoccupied area can be physical or abstract, i.e. “the space between planets” or “a teenager needing their own personal space”. The word can also be used as a verb to refer to a person physically or emotionally distancing themselves from a situation, i.e. “spacing out during lectures”.

Polysemy introduces ambiguity to a retrieval if a query is not given enough context and the system retrieves the undesired version of the word.

- (b) Synonymy refers to different words addressing the same meaning. For example, *colossal*, *giant* and *huge* all describe the size of an object to be very big. Numerous other words also act as synonyms for *big*. Synonymy can lower performance of a retrieval system if it is not aware of the numerous synonyms a query word may have. If a user queries for “big company” but the system only contains documents with the numerous synonyms of *big*, the ranked documents may not be what the user implied.
- (c) Morphology